

How Can AI Augment Access to Justice? Public Defenders' Perspectives on AI Adoption

ANONYMOUS AUTHOR(S)*

Public defenders are asked to do more with less: representing clients deserving of adequate counsel while facing overwhelming caseloads and scarce resources. While artificial intelligence (AI) and large language models (LLMs) are promoted as tools to alleviate this burden, such proposals are detached from the lived realities of public defenders. This study addresses that gap through semi-structured interviews with fourteen practitioners across the United States to examine their experiences with AI, anticipated applications, and ethical concerns. We find that AI adoption is constrained by costs, restrictive office norms, confidentiality risks, and unsatisfactory tool quality. To clarify where AI can and cannot contribute, we propose a task-level map of public defense. Public defenders view AI as most useful for *evidence investigation* to analyze overwhelming amounts of digital records, with narrower roles in *legal research & writing*, and *client communication*. *Courtroom representation* and *defense strategy* are considered least compatible with AI assistance, as they depend on contextual judgment and trust. Public defenders emphasize safeguards for responsible use, including mandatory human verification, limits on overreliance, and the preservation of relational aspect of lawyering. Building on these findings, we outline a research agenda that promotes equitable access to justice by prioritizing open-source models, domain-specific datasets and evaluation, and participatory design that incorporates defenders' perspectives into system development.

CCS Concepts: • **Do Not Use This Code** → **Generate the Correct Terms for Your Paper**; *Generate the Correct Terms for Your Paper*; Generate the Correct Terms for Your Paper; Generate the Correct Terms for Your Paper.

Additional Key Words and Phrases: public defenders, criminal justice, artificial intelligence, large language models, generative AI, professional ethics, AI ethics, qualitative study, participatory design

ACM Reference Format:

Anonymous Author(s). 2024. How Can AI Augment Access to Justice? Public Defenders' Perspectives on AI Adoption. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 23 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

In the United States, public defenders provide representation to defendants who cannot afford private counsel. However, public defenders are systematically overworked and under-resourced [89]. This exacerbates disparities in the justice system: underprivileged groups are disproportionately harmed when they receive inadequate representation [15, 32, 101]. Meanwhile, Artificial Intelligence (AI) and Large Language Models (LLMs) are reshaping work across industries [69, 79, 93]. General-purpose tools, alongside specialized applications, promise to automate or accelerate legal tasks. Scholars have suggested that AI could help close the justice gap, by developing specialized solutions assisting public defenders [72]. However, apart from anecdotal evidence about commercial vendors and some public defender offices experimenting with AI tools, little is known publicly about how public defenders themselves view the use of AI in their work.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2024 Copyright held by the owner/author(s). Publication rights licensed to ACM.

Manuscript submitted to ACM

Manuscript submitted to ACM

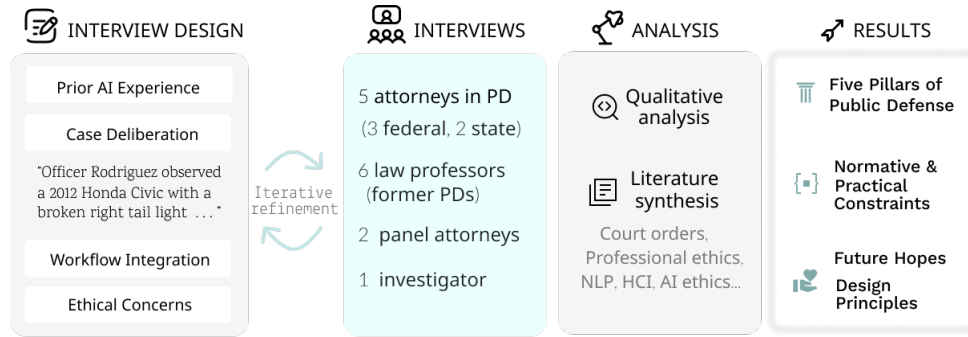


Fig. 1. Overview of our research process: we interviewed 14 experts in public defense on their AI experience, case deliberation, workflow integration, and ethical concerns; applied qualitative analysis and literature synthesis; and distilled insights into five pillars of public defense, normative and practical constraints, and design principles for responsible AI adoption.

To address this gap, we conducted semi-structured interviews with 14 current and former public defenders across the United States. We summarize our methodology in Figure 1. Our study addresses three research questions:

- RQ1: Which areas of public defense work present opportunities or limitations for AI integration?
- RQ2: What technical, organizational, and cultural factors influence the adoption of AI in public defense?
- RQ3: How can we ensure responsible AI use in public defense to uphold professional ethics and quality of work?

To answer these questions, we conduct a qualitative analysis of semi-structured interviews with fourteen participants who have substantial experience in public defense. We find that public defenders see the greatest promise in *evidence investigation*, where AI could transcribe, summarize, and synthesize large volumes of digital records. They assign narrower roles to *legal research and writing*, which requires strict manual verification, and to *client communication*, where AI may assist with clarity. *Courtroom representation* and *defense strategy* are viewed as least compatible with AI assistance, as they rest on contextual judgment and human relationships. Across the board, AI adoption is constrained by four barriers: costs, restrictive office norms, confidentiality risks, and unsatisfactory tool quality. Public defenders also stress safeguards for responsible use, including mandatory human verification, explicit limits on overreliance, and preserving the relational dimensions of lawyering.

This study advances the literature in several ways. It centers the voices of public defenders and identifies the tasks they most want supported by AI, aligning with broader calls to ground AI development in the actual needs of workers [99]. It introduces a task-level map of public defense through five pillars of practice, challenging simplistic claims that generative AI signals attorney incompetence or that every output must be manually verified. Finally, it outlines a forward-looking research agenda for access to justice, highlighting the potential of open-source models to reduce vendor dependence while ensuring confidentiality and cost-efficiency. Building such systems require sustained effort from research communities, including the development of domain-specific datasets, evaluation benchmarks rooted in real public defense practice, and participatory design with public defenders and other stakeholders.

2 Background and Related Work

This section situates our study within the broader legal and technical literature. We outline the structure and challenges of U.S. public defender systems and review the emergence of AI in criminal defense, highlighting both its applications

and the risks. Finally, we draw on prior scholarship examining how professionals adopt AI in high-stakes domains to illuminate the need for empirical research on public defenders.

2.1 Public Defender Systems

The U.S. established public defender systems to ensure legal representation for indigent criminal defendants [111]. After *Powell v. Alabama* sparked the nationwide concern that defendants facing the death penalty received almost no legal representation in 1932 [1, 111], eventually in 1962, the U.S. Supreme Court declared the constitutional right to counsel in *Gideon v. Wainwright* [2]. In response to Gideon’s hand-written petition when he was in jail, the Court declared the state’s obligation to provide indigent (poor) defendants an attorney. Justice Hugo Black argued: “[I]n our adversary system of criminal justice, any person hauled into court, who is too poor to hire a lawyer, cannot be assured a fair trial unless counsel is provided for him. . . . [L]awyers in criminal courts are necessities, not luxuries.”

The constitutional guarantee is not a guarantee of resources, and public defenders in many states struggle: time scarcity, resource limitations, and overwhelming caseloads limit their ability to provide adequate legal representation [84, 89, 103, 112]. Utah’s public defenders were assigned as many as 525 misdemeanor cases per person, and some Louisiana public defenders have caseloads as high as 1,000 clients. This only allows spending a few hours per client [91]. New Mexico Chief Public Defender Bennett Baur was held in contempt of court, after his office failed to appear in five criminal cases due to the caseload strains [68]. Excessive caseloads undermine the quality of representation for indigent defendants [49, 70, 111].

These challenges manifest differently depending on how states structure their public defense systems. In some states, public defender offices employ full-time attorneys and staff dedicated exclusively to representing indigent defendants. In other cases, assigned or contracted counsel consists of private attorneys who take cases through court appointments, government contracts, or nonprofit-managed programs. For example, Washington DC and New Jersey, have relatively well-resourced public defenders’ offices. Michigan, Pennsylvania, and Texas heavily rely on individual counties, without sufficient state funding [43, 81]. To address such systemic resource constraints, our work asks whether AI technologies can help bridge this gap in indigent defense.

2.2 AI Adoptions and Concerns in Criminal Defense

The emergence of LLMs capable of interpreting formal legal language [11] and document summarization [78, 117] has significantly broadened AI’s accessibility and application in legal practice [10]. The capacity of LLMs expanded both the depth and breadth of AI tools in criminal defense [62]. There are two categories of LLM-based tools. **General-purpose AI tools** (e.g., ChatGPT, Claude, Gemini) support broad tasks like research, brainstorming, general information seeking, drafting, and summarization, while **specialized AI tools** (e.g., Westlaw AI, Lexis AI, discovery platforms) are designed for specific legal functions such as case analysis, evidence review, and transcription [28]. These tools promise efficiency but also raise concerns such as hallucinations, bias, and breaches of confidentiality. These concerns are especially acute for public defenders, who serve as the last line of defense for indigent clients and handle their most intimate data, from jail calls to medical records. A single misstep can result in a wrongful conviction.

Hallucinations. As LLMs are trained to predict the most likely next token given the context, they are good at producing plausible-sounding, but not necessarily factually grounded answers [71]. Because the problem stems from the architectural logic of generative models, mitigating hallucinations is one of the most lively research fields [80, 86]. Retrieval-Augmented Generation methods that aim to ground models beyond their internal representations [64, 95]

have shown limited improvements [71]. Despite marketing claims of “hallucination-free AI” [65], independent research finds that Lexis AI and Westlaw AI still produce incorrect information 17–34% of the time [71]. As of September 28, 2025, one tracker has documented 246 incidents of fabricated citations identified by courts, involving 135 attorneys, 103 pro se litigants (individuals representing themselves without attorneys), and 3 judges [47]. Courts have expressed deep frustration with the flood of fabricated cases. As one California judge remarked, “fabricated citations and erroneous statements of law have required this court to spend excessive time,” harming “taxpayers” and “many of whom wait years for a resolution.” [3] Judges nationwide have issued more than 100 orders addressing the problem, ranging from bans on AI-generated filings to mandates requiring disclosure of AI use [53]. Reflecting these concerns, the American Bar Association (ABA)’s Formal Opinion 512 emphasizes that lawyers must be aware of the risks of generative AI, verify all outputs used in legal practice, and maintain supervisory responsibilities over how such tools are used [13].

Bias. Bias in AI outputs have raised concerns in criminal justice for their prevalent use in predictive policing, and criminal risk assessment [16, 17, 31, 56, 105]. LLMs are well-documented to replicate and reinforce the biased present in their training data [29, 45, 50]. For example, vision-language models associate “white” with globalist, “Middle Eastern” with terrorist, and “East Asian” with authoritarian [45]. Another study finds that vision models generate sexualized images 73% of the time for a “17-year-old girl” prompt, but less than 9% for comparable male prompts [113]. Language models shows “dialect prejudice,” such as recommending harsher sentences, including the death penalty, for speakers of African American English [49]. Such biases could distort narratives when LLMs are applied to real-world public defense tasks, such as transcribing jail calls in African American English, automating case intake processes, conducting forensic analysis of devices seized for alleged possession of inappropriate material, or facilitating client communications where dialect, tone, or expression may be unfairly judged.

Confidentiality. LLMs and LLM-integrated services accessed through APIs present inherent security vulnerabilities [33]. The centralized accumulation of vast amounts of data creates opportunities for adversaries to exploit techniques such as model inversion, membership inference, or extraction attacks to obtain sensitive information [37, 110, 114]. Another concern arises from the practice of many LLM providers using user records to retrain or fine-tune their models. For instance, OpenAI discloses that ChatGPT records may be used for further training [85]. Given LLMs’ well-documented tendency to “memorize” data, this raises the risk of inadvertent storage and reproduction of sensitive information from training data [67]. More fundamentally, data preserved by the AI or cloud service providers may fall outside attorney–client privilege or attorney work product, exposing them to disclosure upon civil subpoenas, criminal discovery orders, search warrants, or national security letters [23, 90]. Sam Altman also acknowledged this risk [92]. As courts have treated search histories and emails as discoverable [27, 107], inputs and outputs on the third-party platforms risk being analogized to such correspondence, revealing facts unfavorable to clients or internal strategy-building.

2.3 Professionals’ Perspectives on AI Adoption

AI adoption in high-stakes domains has been widely examined in light of both its potential and its pitfalls. A growing body of research emphasizes the importance of hearing directly from professionals who experience technological needs in practice and whose work is guided by professional ethics [39, 109]. Qualitative interviews have explored physicians [79], home care workers [102], and mental health professionals [69]. By interviewing legal professionals, studies surface the conditions under which AI should provide legal advice to the public [23] and legal professionals’ trust in such systems [58]. There are qualitative studies particularly on public defenders, but they examine contested computational forensic tools [55] or the impact of body-worn camera footage on adjudication [96]. Little is known about

public defenders' broader perceptions of AI technologies, particularly in the wake of LLMs. This study addresses that gap by examining how public defenders view AI's potential to support chronically understaffed offices, the work-intensive tasks most amenable to AI assistance, and the ethical considerations involved in legal representation.

3 Methods

We conduct semi-structured interviews to build new knowledge around the perceptions of public defenders on AI [60]. Our research was approved by our Institutional Review Board.

Expert Recruitment. From March to August 2025, we conducted interviews with 14 individuals, 13 of whom are current or previous public defenders and one is an investigator at a public defender office (*see* Appendix A). We contacted more than ten public defender offices and related non-profits across the United States. We were invited to one federal public defender office's weekly meeting, where we pitched the project and secured 5 participants. However, getting interviews with current public defenders proved challenging: First, public defenders maintain demanding schedules that leave little time for additional commitments. Second, both our institution and public defender offices prohibited providing compensation for interviews as it could constitute lobbying. Through our institutional social media accounts, newsletters, listservs, Slack channels, and personal networks, we widely advertised the study.¹ In addition, we reached out to 150+ attorneys who identified their public defense experience through customized emails and LinkedIn messages. Law professors who are former public defenders, mostly leading the criminal defense clinics, were particularly responsive to our recruitment efforts, resulting in 6 interviews with them. All interviews were conducted via Zoom with participants' consent. We generated automated transcriptions, followed by manual corrections. For interviewees not employed by government, we offered \$50 as compensation for one-hour interviews.

Interview Structure. We conducted semi-structured interviews [5, 8, 54], mostly 1:1, with two sessions involving two participants. The interviews have two parts: (1) **case-based reasoning exercises** and (2) **open-ended questions**. We presented realistic fact patterns (an example can be found in Appendix B) commonly encountered by public defenders, and we observed how participants identified legal issues and come up with defense strategies. This method is widely used to elicit experience-based knowledge among professionals with clinical experience [23, 35, 59, 82]. The open-ended questions evolved iteratively throughout the data collection process in line with the emergent nature of qualitative inquiry, where "preliminary analysis inform[s] data collection." [5, 6, 46, 98]. Since we initially hypothesized that public defenders would use AI mainly for legal research and writing, we prepared demonstrations of tools like ChatGPT and Westlaw AI. However, early interviews showed that research and writing were only part of their work, and participants identify evidence review as more suitable for AI support. This prompted us to adopt a broader view of the work pipeline and develop more detailed questions about workflow, professional identity, and attorney-client relationships. We continued recruitment and data collection until saturation—the point at which additional interviews yielded no novel insights or themes [40]. Appendix C documents the evolution of our interview protocol.

Analysis. We analyzed 12 transcripts, with total word count of 84,908 using abductive coding [23, 104]. Consistent with grounded theory methods, we used emergent themes from data collection to concurrently refine our interview questions and inform our analytical framework. Our findings enter into dialogue with literature from AI integration in legal professions [7, 11, 19, 38, 83] and its normative and practical concerns [23, 34, 63]. Our analysis synthesizes relevant aspects of these fields within the context of our research questions. One author cleaned transcripts to preserve

¹The recruitment website is blinded for reviews.

conversational tone while enhancing readability. Three authors engaged in iterative codebook refinement through multiple team meetings, ensuring that all transcripts were assigned to at least two coders. The codebook was finalized through multiple rounds of meetings. In this process, multiple coders examined all documents and reached consensus on codes through discussion, rendering inter-rater reliability metrics unnecessary [75].

4 Results

Our main results consist of insights on how AI tools can potentially benefit public defenders. All participants have some experience using AI tools, and several attended AI training. However, most participants do not currently use it regularly in their work. P8 and P10 reported their use of specialized tools, such as transcription software for video and audio evidence. P2 and P7 reported deep familiarity with ChatGPT and other tools. All other participants (71%) stated they had not used AI for professional use. We first discuss barriers to wider adoption, followed by a categorization of public defenders' tasks summarizing their workflow and identifying opportunities for AI assistance. We conclude with key ethical considerations participants raised around AI usage.

4.1 Barriers Preventing Adoption

We identify four primary reasons for limited adoption: (1) prohibitive costs especially for self-employed attorneys, (2) office norms and policies, (3) unresolved concerns around confidentiality, and (4) current tools' unsatisfying results.

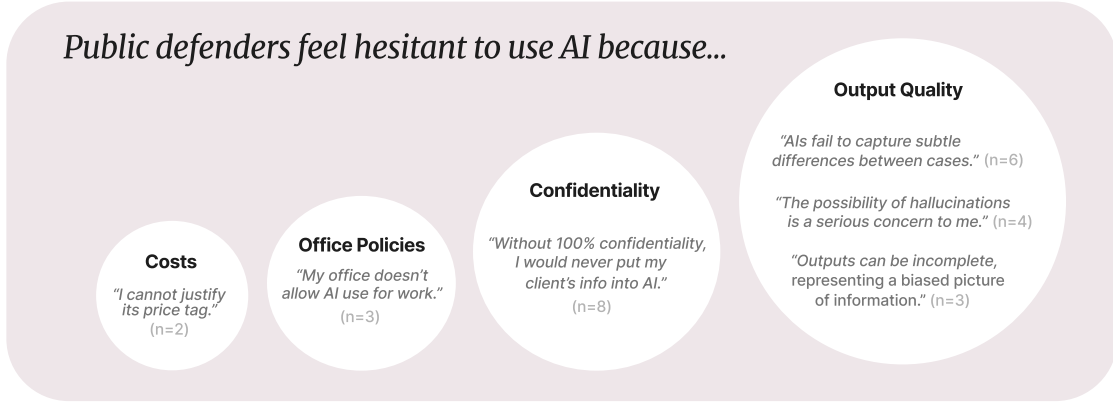


Fig. 2. Main barriers responsible for AI tool adoption in public defense work, grouped into three categories (confidentiality, output quality, and office policies). The darker the color of a box, the more often it has been mentioned.

Costs. Public defender offices and individual attorneys often lack the resources to experiment with or purchase commercial systems. P5 explained that all expenses must pass through county approval, a process that is both slow and uncertain: "It's expensive. I work for a public agency so every kind of new technology needs to be properly tested out and vetted, and the funding has to be approved. So it's a little bit tricky that way." The financial challenge is more acute for private attorneys who take on public defense cases without institutional backing. As P2 shared, "I don't have a Westlaw subscription just because I can't afford it, and I'm not part of an office. The only legal database I have access to is the free one that comes with the Washington State Bar, and that's Fastcase." For these attorneys, even basic legal research requires navigating resource limitations, making AI services inaccessible. P14 explained that public defenders "scrounge for

investigation on [their] own budget,” which falls far short of purchasing the automated synopses of body-worn camera footage that prosecutors have access to.

Office norms and policies. For public defenders employed by government offices, AI use is governed by institutional policies. Some offices approved enterprise-level systems only (“*our offices said, stay away from it outside of Lexus or Westlaw.*” (P1)), whereas in others no formal rules exist, and attorneys are guided only by broad cautionary advice. Private attorneys who take on public defense cases part-time are not bound by office policies and therefore have greater freedom to experiment with AI. Beyond workplace rules, participants noted the court warnings. P2 mentioned, “*I think some judges have issued rulings or standing orders that requires attorneys to disclose if they’ve used any Generative AI.*” but they found such orders “*vague and unenforceable because no one can really agree on a definition of what Generative AI means [such as whether it would include] a simple auto-correct or spelling mistakes.*” Reflecting the growing attentions and concerns around AI, more continuing legal education (CLE) programs (the ongoing professional education required for attorneys to maintain their licenses) offer courses on generative AI but our participant (P7) reported that many trainings are not helpful as many of them provide potential pitfalls, not applicable best practices.

Confidentiality concerns. Confidentiality is the foremost concern for all participants. Participant expressed that they would never submit any client-identifying information into AI systems, unless a system received formal approval from their office or underwent rigorous confidentiality clearance procedures. Participants expressed a strong preference for in-house AI tools that keep data within a “*closed universe.*” As P7 explained, “*I would like a closed universe. I would like to be able to give as much information as possible to the tool to be effective, and confidentiality is what limits me from doing that.*” Similarly, P13 was aware of several law firms purchasing “*closed versions of ChatGPT.*” P5 underscored the challenge, noting that in most applications data “*necessarily has to go off site,*” processed on external servers, and that it “*would be nice if we could do it without taking any of that offsite.*”

Some participants also emphasized that protective orders explicitly restrict the use of AI. P3 and P4 explained that prosecutors provide unredacted sensitive discovery, such as bank records or warrants identifying confidential informants, only under agreements that this material will not be shared outside the office or submitted to AI systems. While these protective orders help public defenders access complete information, they also codify limits that make AI assistance unavailable, even for transcriptions. Not all participants, however, framed confidentiality as an absolute barrier. P2 suggested that concerns may sometimes be selectively applied: “*when people say they’re worried about confidentiality, that is a perfectly valid reason to be skeptical, but then they’ll use Gmail or Google Workspace or Dropbox to host files and scroll to the bottom and click agree.*”

Unsatisfying Output Quality. More than half of participants (n=8) described current AI outputs as unhelpful or unreliable for their work. While tools could return surface-level matches, they frequently missed substantive legal issues. As P1 recalled: “*I was asking about the standard for ordering mistrial based on post-trial conduct, and it gave me a case in which there had been a mistrial in the guy’s first trial, but this decision was about a subsequent trial, so the issue of mistrial wasn’t part of what the case was about.*” Similarly, participants noted that while tools produced technically accurate results for “*black-letter law,*” they lacked sensitivity to factual distinctions and the nuanced case-by-case analysis that characterizes skilled legal reasoning. As P5 explained: “*[i]t feels superficial to me. It feels as if the analysis is made on a very literal-minded level, and it doesn’t seem to really be able to read into those facts what the significance is of them, not what the definition is of the words, but what those facts mean in terms of the weight of the evidence.*”

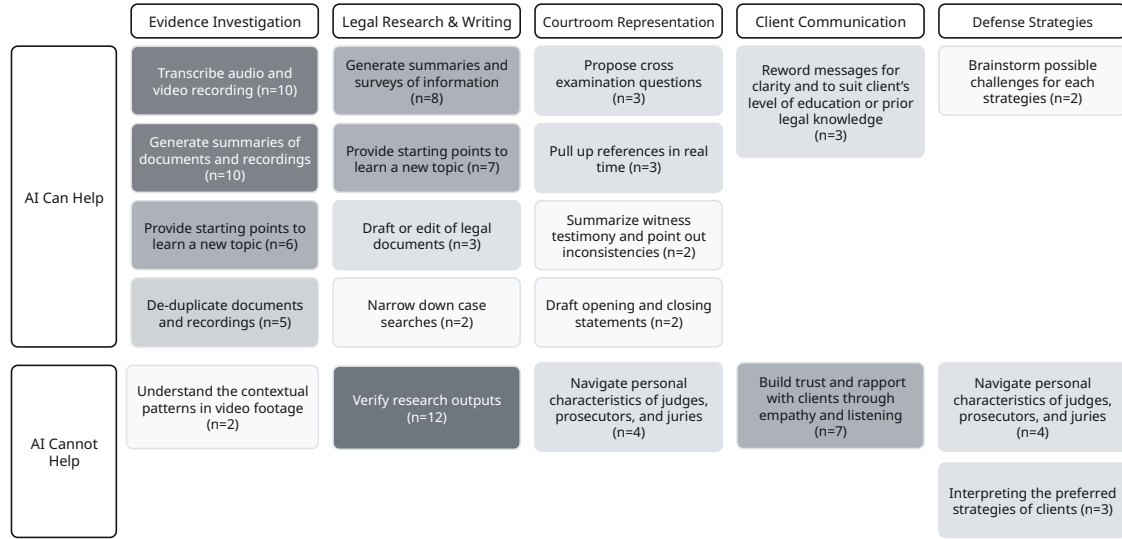


Fig. 3. Ways AI tools can help in public defense work. Over half of the participants are optimistic about using AI tools to help with legal research and writing, as well as evidence investigation.

Others emphasized that beyond obvious citation errors, more subtle hallucinations, such as misleading summaries or inaccurate interpretations, were more difficult to detect and could pose risks if unnoticed. P13 shared that “[my] experience with the Westlaw and Lexisnexis [AI] is they’ll produce text that sounds reasonable, but they often cite to cases that do not stand for the principles that they say the case says. [For example,] if you ask them a question about if a person has an insanity defense, will Missouri recognize temporary insanity, and it might give an answer that says the courts in Missouri are split on whether or not it will be recognized, but this court says it will if these conditions are met, and then they cite to a case. But when you read the case, the case does not say what the AI assistant said it says.” Despite this, some participants believed AI could accelerate certain tasks if used carefully. P2 described their experience: “I asked [ChatGPT] ‘can you point me to the RCW, the Washington State Code,’ and it got it wrong the first time, but with [me] working it and babysitting it, [I] found [my] answer relatively quickly.”

4.2 What AIs Can and Cannot Do: Five Pillars in Public Defense Work

At the outset, we expected legal research and writing to be the main site where AI could play a role. As conversations unfolded, it became clear that public defenders’ work involves many more components, each presenting distinct opportunities and risks for technological support. Through this iterative process, we identified five broad categories that capture their everyday practice: (1) Evidence investigation, (2) Legal research and writing, (3) Courtroom representation, (4) Client communication, and (5) Defense strategy. For each, we distilled specific tasks from participants’ accounts and considered where AI assistance may prove most valuable.

Evidence investigation. Evidence investigation includes reading police reports, watching body camera footage, listening to audio recordings, as well as reviewing all other documents relevant to the case. Most participants noted that this step is **among the most time consuming** in defense work, and could benefit heavily from AI assistance. They

shared that it is often impossible to go through all the materials, due to the sheer volume and time constraints. The prevalence of body-worn cameras and phone calls usually lead to hours and hours of recordings that public defenders need to sift through. The recordings are informationally sparse, but it is difficult to know a priori which portion to watch or listen to. Police reports' descriptions usually provide some pointers, but they might not be comprehensive either and may give a biased representation. Most participants (n=12) mentioned that they are **eager to use AI to help make sense of discovery data**. The most frequently mentioned tasks are transcribing videos or recordings and generating summaries. Defenders can search through transcriptions using keywords or parse them quickly: *"if only we could just send it all out for transcription, that would be a good start because I read a lot faster than I can hear."*(P5) Similarly, summaries help triage whether a document is worth reading or a video is worth watching.

Participants also mentioned other methods of reducing the large volumes of data. Six participants expressed interests in using AI for span extraction, which is identifying key segments in videos, such as asking the AI tools *"Can you give me the timestamp of when he reads the Miranda rights to the defendant"* (P6), *"identify an area within a video [...] let's say warehouse door opens, and you don't want to go through six hours of video to find out when exactly that happens,"* (P2) and *"The crime occurs at [some location] at 4am, and there are some events that happen before that that might be pertinent to the event itself. The prosecutors turn over 12h of video, and probably 10 of those hours have nothing to do with the actual crime [...] AI could be used to identify the key parts of the video."* (P9) Five participants mentioned using AI to de-duplicate or piece together recordings and documents: *"I have a [...] case right now with over a hundred thousand records [...]. And there are so many duplicates. [...] And also be able to identify things that are like it's the same document."*

Some offices already use AI tools for discovery review. Two participants shared that they use commercial AI tools for the tasks mentioned above in their office. However, while participants agreed AI could meaningfully assist with routine evidence review, they also emphasized that the **thought process during evidence investigation** is still guided by humans: *"it's not as if the program is doing the actual work, but it's framing things so that [we] know what point to fast forward to, and that just saved [us] two hours."* (P5) Additionally, they stressed that they would always verify the output of the AI tools. For example, they would search for key-words in the AI-generated transcripts, and would verify the timestamp of matched keywords in the original recording to ensure the transcript is accurate.

Legal research and writing. Legal research typically involves researching case law, statutes, or other legal and scientific materials, as well as drafting legal documents. Many participants (n=8) pointed to AI's usefulness in generating summaries or **surveys of large volumes of legal information**. P1, an experienced research attorney, noted the difficulty of traditional Boolean searches, where *"95% of the cases are going to be bad and 5% good,"* and describing winnowing as *"one of the most difficult tasks in criminal law research."* AI tools can speed up the process of identifying the most relevant information. Similarly, P10 envisioned AI's potential in expansive multi-jurisdictional surveys: *"you can say, how does every State treat guilty except for insanity defenses [...] and it would go survey all the States [...] that could be really persuasive where you could say, our state is the outlier."* These reflections highlight how AI could accelerate tasks that defenders currently experience as laborious and time-consuming. Participants also described turning to AI for quick entry points into new topics. Seven participants said they would use AI tools to quickly orient themselves: P2 reported that they often *"ask [ChatGPT] to generate quizzes to test [their] knowledge and retention,"* and P4 highlighted AI tools as *"a launching point for additional cases."*

Participants also remarked that cases sometimes require research outside the law. As P11 explained: *"If [you] have a scientific area that's relevant in [your] case, [you] need to become a mini-expert about that, and that's very hard for a public defender to do with time constraints."* There are materials which are not indexed by WestLaw, such as finding

relevant academic articles where AI could retrieve relevant documents. For example, P12 noted “*I recently used ChatGPT to find out if there were any law review articles on a certain topic that I was working on [...] and I tried ChatGPT, and it was really helpful because it searched Google*”. In such cases, participants believed AI’s ability to summarize and explain large bodies of information could be especially valuable. Notably, many participants said that legal research & writing was a **relatively small portion** of their jobs. When drafting motions, they usually rely on templates or past motions, circulated in their office, instead of using AI. A few had experimented with AI to generate templates or edit drafts. For example, P11 found AI beneficial for editing writing: “*running your motion or your brief into ChatGPT, and say, edit this for clarity and concision you’re gonna get, It’s going to be an effective use, active tool.*” (P11) But even in those cases, the bulk of the writing remained their own.

Courtroom representation. This category involves presenting arguments, examining witnesses, and persuading judges and juries. Most participants believed that AI has only a limited role in this domain, though some saw potential for assistance in preparation. As P14 explained, “*Your cross-examination in court isn’t spontaneously conceived. If you’re a good lawyer, you are preparing for that in advance, and so that’s one other category that you could consider [AI adoption].*” Others similarly imagined AI **supporting modular and auxiliary tasks** such as summarizing witness testimony, drafting opening or closing statements, or quickly pulling references (P5, P6, P11, P12).

Looking ahead, a few participants speculated about more ambitious possibilities. P7, for instance, suggested that future tools might one day listen to oral arguments and provide real-time suggestions. But participants also highlighted why this is unlikely to be practical in the near term. As P5 emphasized, **courtroom tools would need near-perfect reliability**, since attorneys cannot stop mid-argument to verify an AI’s output. For some, the deeper concern was not technical but professional. P8 argued that oral advocacy is central to the lawyer’s identity: “*That’s where the lawyer comes in.*” Attorneys make emotional arguments and telling stories that resonate with human juries and judges, which, some participants think, AI cannot do. They also cautioned that the use of AI in court would depend on jurisdictional rules, since some prohibit even recordings or electronic devices in the courtroom.

Client communication. This category includes in-person and virtual interactions with clients to explain legal processes, share case updates, and discuss plea options versus trial risks. Ideally, defenders maintain regular communication with clients, though caseload pressures often interfere. Most participants believe AI tools can assist defenders in client communication, but are not able to replace them. Some participants (n=3) already use or are open to using general-purpose AI tools to **reword messages for clarity** or to summarize documents for clients. P7 shared that they would use AI tools to help draft letters and emails to clients. They want to ensure “*[they] do not speak in legalese or convoluted ways, and ChatGPT helps keep [them] in check so that [their] parlance is a little bit more tailored to [their] audience.*” P8 mentioned that his “*clients [...] have not gone that far in school. And so, one thing that you could use AI for is you could write a letter and then ask it, please, make edits to the letter to bring it to a 6th grade educational level.*”

Others imagined AI helping to create visual aids such as flowcharts or decision trees of the legal process that clients could keep as references after the meetings. P8 noted he often uses AI for “*massaging varying kinds of legalese writing into more common parlance description of what’s happening. So you’ve got a very dense paragraph about the grand jury process and cross grand jury notice. And basically how this all works. And then using ChatGPT to boil that down to a more easily communicable idea of: you’re allowed to testify at this proceeding, but we recommend that you don’t.*” Still, nearly all emphasized that **direct communication cannot be delegated entirely**: “*AI cannot substitute having a strong, healthy connection with [our] client and making [our] client feel like [we] care about them, and that [we’re] doing everything [we] can to help them.*” (P11) Moreover, some participants, especially law professors leading criminal defense clinics,

expressed concern that over-reliance on AI could prevent new defenders from developing crucial relationship-building skills. By outsourcing these interactions to AI, students would have fewer chances to develop such skills.

Defense strategy. This category involves deliberate legal and tactical choices to protect clients against criminal charges. Public defenders develop case narratives that aim to exonerate, reduce culpability, or humanize the client, while also deciding whether to file suppression motions, challenge unreliable evidence, or negotiate pleas. A few participants imagined constructive roles of AI in strategizing. P6 suggested that AI might assist with “*brainstorming and also seeing how certain strategies play out... laying out the pros and cons of different options.*” P7 envisioned a tool that could generate strategy suggestions and help identify weaknesses in the prosecution’s case to strengthen negotiations.

However, because strategy blends procedural maneuvering, narrative building, and negotiation, participants also expressed skepticism about AI’s role. For many, strategic judgment was **inseparable from attorneys’ lived experience**: “*I think I have more experience than ChatGPT does in representing clients*”, stated P14. It also involves the **preferences of clients**: “*a big part of a decision about strategy has to come from client, you don’t want to get the answers from AI*” (P14). Both P10 and P14 stated that AI might have seemed useful early in their careers, before accumulated wisdom compressed the need for open-ended brainstorming. For example, drawing on their experience, P14 explained that misdemeanor clients in Washington, D.C. have almost no incentive to enter plea bargains, since the immediate penalties are minor and a guilty plea creates a lasting record, unless multiple convictions accumulate. In addition, P11 flagged the confidentiality concern of sharing the strategy with AI: “*Defense strategy is one of the most privileged pieces of information, and I wouldn’t want to just be putting it all in ChatGPT.*” (P11)

4.3 Necessary Conditions to Foster Responsible Use

4.3.1 Mandatory Human Verification. All participants emphasized the importance of human verification. When using AI tools for legal and scientific research, P2 said that they would use the tools as a way to “*kickstart [their] research*” and find the “*vocabulary necessary to do further research*” with humans still steering the direction of the search. They would also always check the citations and references in AI outputs. Similarly, when using AI tools to assist evidence investigation, participants said that they would use AI for a first pass but still watch the recordings themselves to verify the accuracy of the output and to catch more things that AI tools might have missed, especially as the case progresses (P5, P7, P8). Looking ahead, P11 imagined that there needs to be a regulation in place to ensure that anything generated or modified by AI goes through a check conducted by humans because they think that people who are working under time constraints would be tempted to take AI outputs at face value.

4.3.2 Being Aware of Over-reliance. Participants in teaching positions expressed concerns that over-reliance on AI tools is not helpful to developing essential public defense skills such as efficient information gathering, effective communication, responding to unexpected situations, and decision making under pressure. First, participants have noted that learning legal knowledge and skills from books is very different from engaging in public defense work in real life. For example, P12 noted that “*judges and prosecutors[, being humans,] may not act like what the court opinions make it seem like they do. If AI tools are modeled after what court opinions indicate would happen in a courtroom, it’s [going to] give [students] a false sense of what [they are] actually going to experience.*” They further explained that an essential component of public defense skills is the ability to think on their feet and respond to unexpected situations in the moment. Since AI tools do not have the capability nor are they allowed in the courtroom, public defenders will not be able to perfect everything they write through AI in the courtroom, but have to rely on their own abilities. Moreover, participants stated that even though AI tools can assist legal research, it is still important for the public defenders to

perform their own research outside of the AI generated responses. AI responses may not be comprehensive and may be a biased representation of information. This concern ties into the gilding principle of human verification. As P11 noted, if law school students and early career lawyers become dependent on having AI outputs, *“are we going to create a whole generation of lawyers and students and just people who have lost the ability to think creatively on their own? AI makes mistakes, and it’s important to make sure that our critical thinking stays sharp.”*

4.3.3 Valuing Human Relationships and Judgments. A subset of participants (n=5) emphasized that certain aspects of public defense rely on human relationships and human judgment, and therefore cannot be delegated to AI. Unlike technical tasks, these activities draw on contextual awareness, accumulated experience, and the relational value of time and presence. Participants explained that communication with clients depends not only on words, but also on non-verbal cues such as tone, posture, or hesitation. In the words of P10, *“the clients often give either explicit or more subtle feedback, and the students learn from that. They learned that there was a really awkward pause right there, the client reacted kind of harshly to that, and so they get a sense of maybe don’t say that thing again, or maybe word it in a more careful way.”* Participants stressed that these skills can only be learned from engaging directly with the clients and reflecting on these interactions.

Participants also described courtroom advocacy to be deeply tied to experiential knowledge. Public defenders learn preferences, temperaments, and informal habits of specific judges and prosecutors. As P12 put it, *“it’d be really hard for AI to help [lawyers] figure out what the best approach is in a given case because everything is also so personality specific.”* Unlike human attorneys, AI systems do not have this accumulated background knowledge unless the user explicitly inputs it, and even then participants doubted AI can apply it with the same nuance. Participants further stressed that even if AI could one day observe nonverbal cues and generate contextually appropriate responses, clients would still feel undervalued if their main interactions were with a machine: *“Clients have a right to counsel, not to machines.”* (P14) For many, the time a public defender spends with them is itself a sign of care and commitment. This relational value also extends beyond the case itself: defenders often go above and beyond to support clients’ broader well-being. As P14 explains, *“a good public defender does a lot of social work, [such as] getting their client into the drug program, getting their client mental health services, making sure their client gets enrolled in school, or does their community service to complete their diversion programs.”* These efforts are not only instrumental in improving clients’ lives but also signal to clients that their defenders genuinely care about them as people.

5 Discussion

Our findings show that public defenders view AI as beneficial for the most time-consuming tasks such as evidence review but they are constrained by confidentiality rules, cost barriers, and institutional policies. They also emphasize that strategy-building, courtroom advocacy, and client relationships are inseparable from human judgment, highlighting both the promise and limits of AI in practice. Building on these results, we turn to broader issues for the future: how to segment public defense work in ways that align AI adoption with professional values, how to move beyond a simple human-machine error dichotomy, and how open-source models can offer more secure and equitable alternatives.

5.1 Future of Public Defense Work: Beyond AI Assistance in Legal Research & Writing

Current discussions around AI use in the legal profession focus almost entirely on generative AI for legal research & writing [13, 53], a focus we seek to challenge. We believe the current narrative emerged because of overall early excitement around LLMs [20], and work stating that GPT-4 allegedly acing the bar exam [57]. Not surprisingly, NLP

researchers have worked on AI for legal research [e.g., 73] and writing [e.g., 52]. However, this field remains an open area of research, as problems such as hallucinations still persist. Second, some lawyers' irresponsible use of general-purpose LLMs for motion writing resulted in the occurrences of hallucinated citations [47]. The frustrated warnings of judges and headlines about attorneys being disbarred for using generative AI have attracted global attention that influenced the overall perceptions of using generative AI for legal work (see Section 2.2).

However, legal research & writing are only one part of public defenders' practice. Public defenders spend significant portion of their time distilling large records into concise summaries, classifying materials by relevance, synthesizing insights for strategic decisions, and preparing for cross-examinations. Such time-consuming tasks can be assisted by AI. For example, In video understanding, Video LLMs with advanced functionalities in classification, captioning, summarization, and question answering can support public defenders by enabling faster extraction of insights from complex audiovisual records [115]. Table 1 maps some of these tasks to related streams of AI and NLP research.

Table 1. AI/NLP Technologies for Public Defender Tasks

Task Category	Public Defender Need	Related AI/NLP Research
Making Sense of Large Volumes of Data	Extracting relevant information from massive datasets, described as finding "needles in the haystack." Tasks include summarization, span extraction, classifying data worth examining, identifying relevant documents or passages, and ranking data points by relevance.	Document classification [25], passage assessment [116], learning to rank [44, 66], video span extraction [42], agentic web search [4, 30, 87]
Legal Writing and Translating Legalese	Legal search, Drafting legal documents, translating legalese into accessible language for clients and target audiences.	Legal passage retrieval [73], legal analysis generation [51], translating legalese [12]
AI as a Thought Partner/Educational Tool	Using AI for brainstorming, rapid learning of new topics, feedback on drafts, and preparation for cross-examinations or court hearings. Functions as an intellectual sparring partner for legal strategy development.	AI as thought partner/sparring partner applications [26, 108]

In line with Shao et al. [99], we contend that AI/NLP research should be rooted in workers' needs instead of narrow, profitable tasks. Rather than framing the use of generative AI as a result of the laziness of attorneys, we argue for a proactive approach that incorporates advances in AI and NLP across different segments of public defense work. This requires detailed segmentation of tasks, since the ways of adopting AI, the associated risks, and potential mitigations vary widely. Blanket admonitions such as "attorneys must verify all AI outputs" offer little guidance. While checking citations in a motion might take two hours, verifying whether AI correctly flagged events in 200 hours of video is an entirely different burden. In addition, this segmentation must involve normative judgment on preserving the epistemic agency of human actors [22, 74]. Some components of work are quintessential to public defender's efficacy ("*In the courtroom, that's where attorneys shine*" (P8)), and some are essential to clients ("*Clients have a right to counsel, not to machines*" (P14)). These tasks should not be replaced by AI. Identifying this critical sphere of work is essential both for preserving professional identity and for designing healthier modes of human-AI interaction [24, 97]. For future research, we propose a four-step guide to structuring AI empowerment tailored to the segments of public defense work.

- (1) **Pain Points:** Which segment of work can be automated or augmented by AI?
- (2) **Human Agency:** In each segment, what constitutes high-value work for public defenders' focused attention?
- (3) **Risks:** What risks arise from both the nature of the work and the technologies?
- (4) **Mitigations:** How can those risks be mitigated technically or organizationally?

5.2 Beyond the Human v. Machine Error Dichotomy

AI errors are inevitable just as human errors are. The important question is not whether they occur but which errors can be tolerated under what circumstances and how they can be controlled. Even advanced video LLMs can miss fine-grained details in blurry footage [115], and if AI misses decisive evidence during an initial sweep, public defenders may lose opportunities that could be critical for a case. Biases in many models also risk producing distorted narratives in sensitive contexts [106]. Human verification in such cases would require watching hours of footage in detail, undermining the purpose of AI adoption. With digital evidence expanding exponentially (what was once a single dash cam video has become twelve streams of footage from multiple angles), public defenders face a mounting burden that taxpayer-funded staffing cannot meet. Therefore, this challenge should not be reduced to a binary choice between human and machine fallibility [36, 99]. A more productive approach is to design systems that combine human and AI capacities so that fatigue-driven mistakes are mitigated and public defenders can devote more energy to strategy and client communication. The combination of low tolerance for errors, high human fatigue, and costly verification should not be treated as a dead end but as a research opportunity and design space.

5.3 Risks of Vendor Reliance in Legal Technology

In our interviews, we notice the pronounced presence of third-party vendors in this space. More than half of the participants mentioned specific commercial services or described hearing about contracts that other defender offices had signed. These tools are perceived as more secure than open platforms like ChatGPT, but more costly requiring institutional subscriptions. Indeed, “legal tech AI” has become a lucrative business, with \$2 billion invested in startups in this field [41]. Scholars note that fragmented legal workflows (e.g., video analysis, docket management), combined with demands for specialized security guardrails, have enabled startups to secure contracts with large law firms [76]. Consolidation between vendors and big law is further fueled by the scarcity of legal data for training models. Although initiatives such as Free Law Project [94], Caselaw Access Project [61], Pile of Law [48], and Cambridge Law Corpus [88] have sought to democratize access, the core repositories of U.S. law remain behind costly paywalls. Startups therefore have strong incentives to partner with established firms to obtain domain-specific data [76].

This vendor reliance phenomenon raises concern of deepening the justice gap by layering expensive AI services on top of an already proprietary information market, leaving under-resourced actors like public defenders further behind [18]. It is also important to note that the promise of confidentiality is fragile. Vendors may promise encryption, non-retention, and local storage options. However, if these services rely on proprietary back-end models hosted by corporations such as Microsoft or OpenAI, it is unclear whether these systems can be considered a closed universe insulated from outside access. The integration of LLMs inevitably invites the risks of unintentional leaks or adversarial attacks [33]. The inputs and outputs handled by third-party vendors could fall outside of privileged information, subject to mandatory disclosure [23, 90, 92]. More fundamentally, there is no public evaluation benchmarks that allow public defender offices to auditing third-party models whether they live up to their promised performance and confidentiality. Unchecked vendor dominance risks turning AI innovation into another barrier to justice rather than a bridge toward it.

5.4 Pathways for Open-Source AI in Public Defense

Concerns around vendor dependence make a strong case for utilizing open-source models [18, 93]. The most important advantage of the open-source models is that attorneys can run models on-device, in controlled local environments. It eliminates the risk of sensitive evidence leaving the public defender’s office and the inputs and outputs of the systems

are likely to be protected as attorneys' work product. In addition, open models allow more rooms technical adjustments, such as citation retrieval systems, that make outputs more trustworthy [18]. In the past, local models were considered inferior in performance compared to proprietary cloud-based systems, but this gap is narrowing quickly. OpenAI's GPT-OSS-20B, with 20 billion parameters, or vision-language models such as Qwen2.5-VL [14] can now run on common laptops. As model optimization continues, we expect this trend to accelerate. Although local deployment requires upfront investment in hardware such as GPUs, it eliminates recurring expenses like subscription fees or API usage costs. The AI-desired tasks public defenders identified in our study do not require the most advanced version of LLMs. What they need are practical, accessible tools that support high-volume, rule-based tasks like document review, triage, and structure evidence summarization and extraction.

To make open-source LLMs practically adaptable in public defense, two elements are essential: **in-house technical expertise** within public defender offices and **domain-specific LLM research** that fine-tunes models for the segmented tasks we discuss in Section 5.1. For the first, states should consolidate resources to provide organizational support for public defenders. Public defenders are known to thrive in teams, having access to institutional support for training, investigation, and social work [9, 100]. AI adoption will add new demands around data practices and the accumulation of institutional knowledge at scale. For the second, Perron et al. [93] shows that local LLMs achieves near-human accuracy in classifying and extracting substance problems in child welfare investigations. However, much work remains to bring such advances into public defense. High-quality datasets are needed, spanning legal texts, client communications, and long-form video data such as body-worn camera footage. Robust evaluation benchmarks are also essential, not only to measure text-level accuracy [71] but also to assess confidentiality, bias, and utility across various forms of data and contexts. We highlight the importance of participatory design to leverage public defenders' expertise and needs [77, 99]. For example, research shows how adaptive vision-language systems in healthcare can dynamically integrate clinician feedback through reinforcement learning [21]. Together, these efforts would allow open-source models to be tailored to the realities of public defense, delivering tools that are both reliable and accessible.

6 Conclusion

Public defenders carry a constitutional mandate under conditions of scarcity. Through qualitative interviews with current and former public defenders, we examine whether AI can relieve parts of the burden. By categorizing public defense work into five pillars, we find that AI assistance is perceived to be most beneficial for time consuming evidence investigation tasks. We also document responsible AI use for public defenders, which include human verification of AI output, being aware of risks around over-reliance, and acknowledging that public defender work to some extent relies on human relationships and human judgment, which is hard for AI to replicate. We envision a future in which NLP and AI are responsibly integrated into public institutions, strengthening equitable access to essential services and supporting the people who deliver them. While there are no simple solutions, several priorities stand out. First, move beyond the human-machine dichotomy and design interfaces that make humans and AI more effective together. Second, direct research toward tasks central to evidence investigation, such as span extraction and video summarization, which defenders identified as most practical. Third, invest in public defender-specific datasets and benchmarks to foster research and enable auditing of both open- and closed-source systems, including commercial tools. Finally, embrace locally deployable models that lower costs and safeguard confidentiality.

References

- [1] 1932. *Powell v. Alabama*. 287 U.S. 45.

- [2] 1963. Gideon v. Wainwright. 372 U.S. 335.
- [3] 2025. Noland v. Land of the Free, L.P. <https://www4.courts.ca.gov/opinions/documents/B331918.PDF>. No. B331918, 2d Dist. Div. 3 (Cal. Ct. App.).
- [4] Deepak Bhaskar Acharya, Karthigeyan Kuppan, and B Divya. 2025. Agentic ai: Autonomous intelligence for complex goals—a comprehensive survey. *IEEe Access* (2025).
- [5] Omolola A Adeoye-Olatunde and Nicole L Olenik. 2021. Research and scholarly methods: Semi-structured interviews. *Journal of the american college of clinical pharmacy* 4, 10 (2021), 1358–1367.
- [6] Amanda Agan, Matthew Freedman, and Emily Owens. 2021. Is your lawyer a lemon? Incentives and selection in the public provision of criminal defense. *Review of Economics and Statistics* 103, 2 (2021), 294–309.
- [7] Allan Almeida, Shubh Shah, Neer Bhavesh Khandor, and Neha Agarwal. 2024. LEGAL AI: A Tool to help Lawyers Optimize their Performance. In *2024 International Conference on Electrical Electronics and Computing Technologies (ICEECT)*, Vol. 1. IEEE, 1–8.
- [8] Claire Anderson. 2010. Presenting and evaluating qualitative research. *American journal of pharmaceutical education* 74, 8 (2010), 141.
- [9] James M Anderson and Paul Heaton. 2012. How much difference does the lawyer make: The effect of defense counsel on murder case outcomes. *Yale LJ* 122 (2012), 154.
- [10] Farid Ariai, Joel Mackenzie, and Gianluca Demartini. 2024. Natural language processing for the legal domain: A survey of tasks, datasets, models, and challenges. *arXiv preprint arXiv:2410.21306* (2024).
- [11] Elliott Ash, Aniket Kesari, Suresh Naidu, Lena Song, and Dominik Stambach. 2024. Translating Legalese: Enhancing Public Understanding of Court Opinions with Legal Summarizers. In *Proceedings of the 2024 Symposium on Computer Science and Law* (Boston, MA, USA) (CSLAW '24). Association for Computing Machinery, New York, NY, USA, 136–157. <https://doi.org/10.1145/3614407.3643700>
- [12] Elliott Ash, Aniket Kesari, Suresh Naidu, Lena Song, and Dominik Stambach. 2024. Translating Legalese: Enhancing Public Understanding of Court Opinions with Legal Summarizers. In *Proceedings of the 2024 Symposium on Computer Science and Law* (Boston, MA, USA) (CSLAW '24). Association for Computing Machinery, New York, NY, USA, 136–157. <https://doi.org/10.1145/3614407.3643700>
- [13] The American Bar Association. 2025. ABA Formal Opinion 512: The Paradigm for Generative AI in Legal Practice. https://www.americanbar.org/content/dam/aba/administrative/professional_responsibility/ethics-opinions/aba-formal-opinion-512.pdf. (2025).
- [14] Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibao Song, Kai Dang, Peng Wang, Shijie Wang, Jun Tang, Humen Zhong, Yuanzhi Zhu, Mingkun Yang, Zhaohai Li, Jianqiang Wan, Pengfei Wang, Wei Ding, Zheren Fu, Yiheng Xu, Jiabo Ye, Xi Zhang, Tianbao Xie, Zesen Cheng, Hang Zhang, Zhibo Yang, Haiyang Xu, and Junyang Lin. 2025. Qwen2.5-VL Technical Report. *arXiv preprint arXiv:2502.13923* (2025).
- [15] Marea Beeman and Claire Buetow. 2023. *Gideon at 60: A Snapshot of State Public Defense Systems and Paths to System Reform*. Technical Report NIJ Report 307325. National Institute of Justice, U.S. Department of Justice. <https://www.ojp.gov/pdffiles1/nij/307325.pdf>
- [16] Richard Berk. 2017. An impact assessment of machine learning risk forecasts on parole board decisions and recidivism. *Journal of Experimental Criminology* 13, 2 (2017), 193–216.
- [17] Richard Berk, Drougas Berk, and D Drougas. 2019. *Machine learning risk assessments in criminal justice settings*. Springer.
- [18] Rohan Bhambhoria, Samuel Dahan, Jonathan Li, and Xiaodan Zhu. 2024. Evaluating ai for law: Bridging the gap with open-source solutions. *arXiv preprint arXiv:2404.12349* (2024).
- [19] Paheli Bhattacharya, Soham Poddar, Koustav Rudra, Kripabandhu Ghosh, and Saptarshi Ghosh. 2021. Incorporating domain knowledge for extractive summarization of legal case documents. In *Proceedings of the eighteenth international conference on artificial intelligence and law*. 22–31.
- [20] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, Chris Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. Language Models are Few-Shot Learners. In *Advances in Neural Information Processing Systems*, H. Larochelle, M. Ranzato, R. Hadsell, M.F. Balcan, and H. Lin (Eds.), Vol. 33. Curran Associates, Inc., 1877–1901. https://proceedings.neurips.cc/paper_files/paper/2020/file/1457c0d6bfc4967418bfb8ac142f64a-Paper.pdf
- [21] Yiming Cao, Zhen Li, Lizhen Cui, and Chunyan Miao. 2025. Adaptive Human-LLMs Interaction Collaboration: Reinforcement Learning driven Vision-Language Models for Medical Report Generation. In *Proceedings of the Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA '25)*. Association for Computing Machinery, New York, NY, USA, Article 62, 6 pages. <https://doi.org/10.1145/3706599.3719852>
- [22] Inyoung Cheong. 2025. Epistemic and Emotional Harms of Generative AI: Towards Human-Centered First Amendment. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5435335.
- [23] Inyoung Cheong, King Xia, KJ Kevin Feng, Quan Ze Chen, and Amy X Zhang. 2024. (A) I am not a lawyer, but...: engaging legal experts towards responsible LLM policies for legal advice. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*. 2454–2469.
- [24] Mark Coeckelbergh. 2015. The tragedy of the master: automation, vulnerability, and distance. *Ethics and Information Technology* 17, 3 (2015), 219–229.
- [25] Aaron M Cohen, William R Hersh, Kim Peterson, and Po-Yin Yen. 2006. Reducing workload in systematic review preparation using automated citation classification. *Journal of the American Medical Informatics Association* 13, 2 (2006), 206–219.
- [26] Katherine M Collins, Ilia Sucholutsky, Umang Bhatt, Kartik Chandra, Lionel Wong, Mina Lee, Cedegao E Zhang, Tan Zhi-Xuan, Mark Ho, Vikash Mansinghka, et al. 2024. Building machines that learn and think with people. *Nature human behaviour* 8, 10 (2024), 1851–1863.
- [27] 2023) Colorado v. Seymour. [n. d.]. 536 P.3d 1260 (Colo.).

- [28] National Best Practices Commission. 2025. Integrating AI: Guidance and Policies for Prosecutors. <https://pceinc.org/wp-content/uploads/2025/01/20250125-Integrating-AI-A-Guide-for-Prosecutors.pdf>
- [29] Sunhao Dai, Chen Xu, Shicheng Xu, Liang Pang, Zhenhua Dong, and Jun Xu. 2024. Bias and unfairness in information retrieval systems: New challenges in the llm era. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. 6437–6447.
- [30] Xiang Deng, Yu Gu, Boyuan Zheng, Shijie Chen, Samuel Stevens, Boshi Wang, Huan Sun, and Yu Su. 2023. Mind2Web: Towards a Generalist Agent for the Web. *arXiv:2306.06070* [cs.CL] <https://arxiv.org/abs/2306.06070>
- [31] David C. Donald. 2025. Bias in AI Large Language Models: Risks and Remedies. <https://www.americanbar.org/groups/gpsolo/resources/magazine/2025-mar-apr/bias-ai-large-language-models-risks-remedies/>.
- [32] John J. Donohue and Eric A. Baldwin. 2024. The Crisis of America's Public Defenders. *Project Syndicate* (8 Oct. 2024). <https://www.project-syndicate.org/commentary/new-public-defender-model-could-transform-us-criminal-justice-by-john-donohue-and-eric-a-baldwin-2024-10>
- [33] Jonathan Evertz, Merlin Chlosta, Lea Schönherr, and Thorsten Eisenhofer. 2024. Whispers in the machine: Confidentiality in llm-integrated systems. *arXiv preprint arXiv:2402.06922* (2024).
- [34] Frank Fagan. 2020. Natural language processing for lawyers and judges. *Mich. L. Rev.* 119 (2020), 1399.
- [35] KJ Feng, Quan Ze Chen, Inyoung Cheong, King Xia, and Amy X Zhang. 2023. Case repositories: Towards case-based reasoning for ai alignment. *arXiv preprint arXiv:2311.10934* (2023).
- [36] Kevin Feng, David McDonald, and Amy Zhang. 2025. Levels of Autonomy for AI Agents Sébastien A. Krier using Midjourney 6.1. *Artificial Intelligence* (2025).
- [37] Georgios Feretakis, Konstantinos Papaspyridis, Aris Gkoulalas-Divanis, and Vassilios S Verykios. 2024. Privacy-preserving techniques in generative ai and large language models: a narrative review. *Information* 15, 11 (2024), 697.
- [38] Jens Frankenreiter and Julian Nyarko. 2022. Natural language processing in legal tech. *Legal Tech and the Future of Civil Justice (David Engstrom ed.) Forthcoming* (2022).
- [39] Yoko E Fukumura, Julie McLaughlin Gray, Gale M Lucas, Burcin Becerik-Gerber, and Shawn C Roll. 2021. Worker perspectives on incorporating artificial intelligence into office workspaces: Implications for the future of office work. *International Journal of Environmental Research and Public Health* 18, 4 (2021), 1690.
- [40] Barney G Glaser and Anselm L Strauss. 2017. Theoretical sampling. In *Sociological methods*. Routledge, 105–114.
- [41] Joanna Glasner. 2026. Legal Tech Startup Investment Is Riding High, Thanks To AI Boost. <https://news.crunchbase.com/venture/legal-tech-startup-investment-ai-clio-harvey/>. *Crunchbase News* (February 2026).
- [42] Kristen Grauman, Andrew Westbury, Eugene Byrne, Zachary Chavis, Antonino Furnari, Rohit Girdhar, Jackson Hamburger, Hao Jiang, Miao Liu, Xingyu Liu, Miguel Martin, Tushar Nagarajan, Ilija Radosavovic, Santhosh Kumar Ramakrishnan, Fiona Ryan, Jayant Sharma, Michael Wray, Mengmeng Xu, Eric Zhongcong Xu, Chen Zhao, Siddhant Bansal, Dhruv Batra, Vincent Cartillier, Sean Crane, Tien Do, Morrie Doulaty, Akshay Erapalli, Christoph Feichtenhofer, Adriano Fragomeni, Qichen Fu, Abrahm Gebreselasie, Cristina Gonzalez, James Hillis, Xuhua Huang, Yifei Huang, Wenqi Jia, Weslie Khoo, Jachym Kolar, Satwik Kottur, Anurag Kumar, Federico Landini, Chao Li, Yanghao Li, Zhenqiang Li, Karttikeya Mangalam, Raghava Modhugu, Jonathan Munro, Tullie Murrell, Takumi Nishiyasu, Will Price, Paola Ruiz Puentes, Merey Ramazanov, Leda Sari, Kiran Somasundaram, Audrey Southerland, Yusuke Sugano, Ruijie Tao, Minh Vo, Yuchen Wang, Xindi Wu, Takuma Yagi, Ziwei Zhao, Yunyi Zhu, Pablo Arbelaez, David Crandall, Dima Damen, Giovanna Maria Farinella, Christian Fuegen, Bernard Ghanem, Vamsi Krishna Ithapu, C. V. Jawahar, Hanbyul Joo, Kris Kitani, Haizhou Li, Richard Newcombe, Aude Oliva, Hyun Soo Park, James M. Rehg, Yoichi Sato, Jianbo Shi, Mike Zheng Shou, Antonio Torralba, Lorenzo Torresani, Mingfei Yan, and Jitendra Malik. 2022. Ego4D: Around the World in 3,000 Hours of Egocentric Video. *arXiv:2110.07058* [cs.CV] <https://arxiv.org/abs/2110.07058>
- [43] Jaelyn Grisso. 2025. Washington among the bottom of states in public defense funding. <https://www.cascadepbs.org/investigations/2025/03/washington-among-bottom-states-public-defense-funding/>
- [44] Jiafeng Guo, Yixing Fan, Liang Pang, Liu Yang, Qingyao Ai, Hamed Zamani, Chen Wu, W. Bruce Croft, and Xueqi Cheng. 2019. A Deep Look into Neural Ranking Models for Information Retrieval. *arXiv:1903.06902* [cs.IR] <https://arxiv.org/abs/1903.06902>
- [45] Kimia Hamidieh, Haoran Zhang, Walter Gerych, Thomas Hartvigsen, and Marzyeh Ghassemi. 2024. Identifying implicit social biases in vision-language models. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, Vol. 7. 547–561.
- [46] David Harper and Andrew R Thompson. 2011. *Qualitative research methods in mental health and psychotherapy: A guide for students and practitioners*. John Wiley & Sons.
- [47] Peter Henderson. 2025. AI Law Tracker – Fake Citations in Court (AI Hallucinations). <https://www.polarislab.org/ai-law-tracker.html>.
- [48] Peter Henderson, Mark Krass, Lucia Zheng, Neel Guha, Christopher D Manning, Dan Jurafsky, and Daniel Ho. 2022. Pile of law: Learning responsible data filtering from the law and a 256gb open-source legal dataset. *Advances in Neural Information Processing Systems* 35 (2022), 29217–29234.
- [49] Morris B Hoffman, Paul H Rubin, and Joanna M Shepherd. 2005. An empirical study of public defender effectiveness: Self-selection by the marginally indigent. *Ohio St. J. Crim. L.* 3 (2005), 223.
- [50] Rachel Hong, Jevan Hutson, William Agnew, Imaad Huda, Tadayoshi Kohno, and Jamie Morgenstern. 2025. A Common Pool of Privacy Problems: Legal and Technical Lessons from a Large-Scale Web-Scraped Machine Learning Dataset. *arXiv preprint arXiv:2506.17185* (2025).
- [51] Abe Bohan Hou, Orion Weller, Guanghui Qin, Eugene Yang, Dawn Lawrie, Nils Holzenberger, Andrew Blair-Stanek, and Benjamin Van Durme. 2024. CLERC: A Dataset for Legal Case Retrieval and Retrieval-Augmented Analysis Generation. *arXiv:2406.17186* [cs.CL] <https://arxiv.org/abs/2406.17186>

- [52] Abe Bohan Hou, Orion Weller, Guanghui Qin, Eugene Yang, Dawn Lawrie, Nils Holzenberger, Andrew Blair-Stanek, and Benjamin Van Durme. 2025. CLERC: A Dataset for U. S. Legal Case Retrieval and Retrieval-Augmented Analysis Generation. In *Findings of the Association for Computational Linguistics: NAACL 2025*, Luis Chiruzzo, Alan Ritter, and Lu Wang (Eds.). Association for Computational Linguistics, Albuquerque, New Mexico, 7898–7913. <https://doi.org/10.18653/v1/2025.findings-naacl.441>
- [53] Responsible AI in Legal Service. 2025. Analysis of AI Use in Courts. <https://rails.legal/resources/resource-ai-orders/>.
- [54] Shazia Jamshed. 2014. Qualitative research method-interviewing and observation. *Journal of basic and clinical pharmacy* 5, 4 (2014), 87.
- [55] Angela Jin and Niloufar Salehi. 2024. (Beyond) Reasonable Doubt: Challenges that Public Defenders Face in Scrutinizing AI in Court. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems* (Honolulu, HI, USA) (CHI '24). Association for Computing Machinery, New York, NY, USA, Article 467, 19 pages. <https://doi.org/10.1145/3613904.3641902>
- [56] Surya Mattu Julia Angwin, Jeff Larson, Lauren Kirchner, and ProPublica. 2016. Machine Bias: There's software used across the country to predict future criminals. And it's biased against blacks. <https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing>
- [57] Daniel Martin Katz, Michael James Bommarito, Shang Gao, and Pablo Arredondo. 2024. Gpt-4 passes the bar exam. *Philosophical Transactions of the Royal Society A* 382, 2270 (2024), 20230254.
- [58] Ryan Kennedy, Lydia Tiede, Amanda Austin, and Kenzy Ismael. 2025. Law Enforcement and Legal Professionals' Trust in Algorithms. *Journal of Law & Empirical Analysis* 2, 1 (2025), 77–96.
- [59] Gary A. Klein, Roberta Calderwood, and Donald G. MacGregor. 1989. Critical decision method for eliciting knowledge. *IEEE Trans. Syst. Man Cybern.* 19 (1989), 462–472. <https://api.semanticscholar.org/CorpusID:28418547>
- [60] Steinar Kvale. 1996. *InterViews: an introduction to qualitative research interviewing*. Sage.
- [61] Harvard Library Innovation Lab. n.a.. Caselaw Access Project. <https://case.law/>.
- [62] UC Berkeley Law. 2025. AI for Defense Lawyers. <https://www.law.berkeley.edu/research/criminal-law-and-justice-center/our-work/ai-for-public-defenders/>
- [63] Michael Legg and Felicity Bell. 2019. Artificial intelligence and the legal profession: Becoming the AI-enhanced lawyer. *U. Tas. L. Rev.* 38 (2019), 34.
- [64] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, et al. 2020. Retrieval-augmented generation for knowledge-intensive nlp tasks. *Advances in neural information processing systems* 33 (2020), 9459–9474.
- [65] LexisNexis. 2023. LexisNexis Launches Lexis+ AI, a Generative AI Solution with Hallucination-Free Linked Legal Citations. <https://www.lexisnexis.com/community/pressroom/b/news/posts/lexisnexis-launches-lexis-ai-a-generative-ai-solution-with-hallucination-free-linked-legal-citations>.
- [66] Tie-Yan Liu. 2009. Learning to Rank for Information Retrieval. *Foundations and Trends® in Information Retrieval* 3, 3 (2009), 225–331. <https://doi.org/10.1561/15000000016>
- [67] Zhe Liu, Xuedong Zhang, and Fuchun Peng. 2023. Mitigating unintended memorization in language models via alternating teaching. In *ICASSP 2023-2023 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 1–5.
- [68] Andy Lyman. 2016. Chief public defender held in contempt after turning down cases, says office can't afford it. <https://nmpoliticalreport.com/2016/12/02/chief-public-defender-held-in-contempt-after-turning-down-cases-says-office-cant-afford-it/>
- [69] Yingzhuo Ma, Yi Zeng, Tong Liu, Ruoshan Sun, Mingzhao Xiao, and Jun Wang. 2024. Integrating large language models in mental health practice: a qualitative descriptive study based on expert interviews. *Frontiers in Public Health* 12 (2024), 1475867.
- [70] Alma Magaña. 2023. Public Defenders as Gatekeepers of Freedom. *UCLA L. Rev.* 70 (2023), 978.
- [71] Varun Magesh, Faiz Surani, Matthew Dahl, Mirac Suzgun, Christopher D Manning, and Daniel E Ho. 2025. Hallucination-Free? Assessing the Reliability of Leading AI Legal Research Tools. *Journal of Empirical Legal Studies* 22, 2 (2025), 216–242.
- [72] Robert Mahari, Dominik Stambach, Elliott Ash, and Alex Pentland. 2023. The Law and NLP: Bridging Disciplinary Disconnects. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, Singapore, 3445–3454. <https://doi.org/10.18653/v1/2023.findings-emnlp.224>
- [73] Robert Mahari, Dominik Stambach, Elliott Ash, and Alex Pentland. 2024. LePaRD: A Large-Scale Dataset of Judicial Citations to Precedent. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Lun-Wei Ku, Andre Martins, and Vivek Srikumar (Eds.). Association for Computational Linguistics, Bangkok, Thailand, 9863–9877. <https://doi.org/10.18653/v1/2024.acl-long.532>
- [74] Emmie Malone, Saleh Afroogh, Jason D'Cruz, and Kush R Varshney. 2025. When trust is zero sum: automation's threat to epistemic agency. *Ethics and Information Technology* 27, 2 (2025), 29.
- [75] Nora McDonald, Sarita Schoenebeck, and Andrea Forte. 2019. Reliability and inter-rater reliability in qualitative research: Norms and guidelines for CSCW and HCI practice. *Proceedings of the ACM on human-computer interaction* 3, CSCW (2019), 1–23.
- [76] Ankit Tandon Megan Ma, Aparna Sinha and Jennifer Richards. 2024. Generative AI Legal Landscape 2024.
- [77] Kazuhiko Momose, Rahul Mehta, Josias Moukpe, Troy R Weekes, and Thomas C Eskridge. 2025. Human-AI Teamwork Interface Design Using Patterns of Interactions. *International Journal of Human-Computer Interaction* 41, 11 (2025), 7112–7134.
- [78] Gianluca Moro, Nicola Piscaglia, Luca Ragazzi, and Paolo Italiani. 2024. Multi-language transfer learning for low-resource legal case summarization. *Artificial Intelligence and Law* 32, 4 (2024), 1111–1139.
- [79] Lina Mosch, Daniel Fürstenau, Jenny Brandt, Jasper Wagnitz, Sophie AI Klopfenstein, Akira-Sebastian Poncette, and Felix Balzer. 2022. The medical profession transformed by artificial intelligence: Qualitative study. *Digital Health* 8 (2022), 20552076221143903.

- [80] Conor Murray. 2023. Why AI 'Hallucinations' Are Worse Than Ever. <https://www.wsj.com/us-news/judge-sanctions-lawyers-who-filed-fake-chatgpt-legal-research-9ebad8f9>. *Forbes* (June 2023).
- [81] George Naufal. 2024. Justice Beyond the Cities: The State Rural Public Defense in Texas. https://ppri.tamu.edu/wp-content/uploads/2024/09/20240529_Rural-Indigent-Defense-in-Texas.pdf A Report for Texas Indigent Defense Commission.
- [82] Maria Skyvell Nilsson and Ewa Pilhammar. 2009. Professional approaches in clinical judgements among senior and junior doctors: implications for medical education. *BMC Medical Education* 9 (2009), 25 – 25. <https://api.semanticscholar.org/CorpusID:12420876>
- [83] Maheshwaran Nithya, S Harini, S Kavyadharshini, and K Srinidhi. 2024. AI-Driven Legal Automation to Enhance Legal Processes with Natural Language Processing. In *2024 International Conference on IoT Based Control Networks and Intelligent Systems (ICICNIS)*. IEEE, 1246–1253.
- [84] Charles J Ogletree Jr. 1992. Beyond justifications: Seeking motivations to sustain public defenders. *Harv. L. Rev.* 106 (1992), 1239.
- [85] OpenAI. 2024. Consumer Privacy at OpenAI. <https://openai.com/consumer-privacy/>.
- [86] OpenAI. 2025. OpenAI o3 and o4-mini System Card. <https://cdn.openai.com/pdf/2221c875-02dc-4789-800b-e7758f3722c1/o3-and-o4-mini-system-card.pdf>.
- [87] OpenAI. 2025. *Operator System Card*. System Card. OpenAI. https://cdn.openai.com/operator_system_card.pdf
- [88] Andreas Östling, Holli Sargeant, Huiyuan Xie, Ludwig Bull, Alexander Terenin, Leif Jonsson, Måns Magnusson, and Felix Steffek. 2023. The Cambridge law corpus: A dataset for legal AI research. *Advances in Neural Information Processing Systems* 36 (2023), 41355–41385.
- [89] Nicholas M. Pace, Malia N. Brink, Cynthia G. Lee, and Stephen F. Hanlon. 2023. *National Public Defense Workload Study*. Technical Report. RAND Corporation.
- [90] Adam Paine and Robert M. Travisano. 2025. Discovery Pitfalls in the Age of AI. <https://techcrunch.com/2025/07/25/sam-altman-warns-theres-no-legal-confidentiality-when-using-chatgpt-as-a-therapist/>. *Epstein Becker Green* (September 2025).
- [91] Jule Pattison-Gordon. 2025. Staffing and Funding Shortages Persist for Public Defenders. <https://www.governing.com/workforce/staffing-and-funding-shortages-persist-for-public-defenders>
- [92] Sarah Perez. 2025. Sam Altman warns there's no legal confidentiality when using ChatGPT as a therapist. <https://techcrunch.com/2025/07/25/sam-altman-warns-theres-no-legal-confidentiality-when-using-chatgpt-as-a-therapist/>. *TechCrunch* (January 2025).
- [93] Brian E Perron, Hui Luan, Bryan G Victor, Oliver Hiltz-Perron, and Joseph Ryan. 2025. Moving beyond ChatGPT: Local large language models (LLMs) and the secure analysis of confidential unstructured text data in social work research. *Research on Social Work Practice* 35, 6 (2025), 695–710.
- [94] Free Law Project. n.a.. Court Listner. <https://www.courtlistner.com/>.
- [95] Keshav Rangan and Yiqiao Yin. 2024. A fine-tuning enhanced RAG system with quantized influence measure as AI judge. *Scientific Reports* 14, 1 (2024), 27446.
- [96] O Nicholas Robertson, John D McCluskey, Shakierah S Smith, and Craig D Uchida. 2024. Body cameras and adjudication: Views of prosecutors and public defenders. *Criminal justice review* 49, 1 (2024), 15–29.
- [97] Advait Sarkar. 2023. Enough with “human-AI collaboration”. In *Extended abstracts of the 2023 CHI conference on human factors in computing systems*. 1–8.
- [98] Jonas Schulze and Felix Pinkow. 2020. Leadership for Organisational Adaptability: How Enabling Leaders Create Adaptive Space. *Administrative Sciences* 10 (2020), 37. <https://api.semanticscholar.org/CorpusID:225684128>
- [99] Yijia Shao, Humishka Zope, Yucheng Jiang, Jiaxin Pei, David Nguyen, Erik Brynjolfsson, and Diyi Yang. 2025. Future of Work with AI Agents: Auditing Automation and Augmentation Potential across the US Workforce. *arXiv preprint arXiv:2506.06576* (2025).
- [100] Yotam Shem-Tov. 2022. Make or buy? The provision of indigent defense services in the United States. *Review of Economics and Statistics* 104, 4 (2022), 819–827.
- [101] Mary C. Slosar. 2022. *The Justice Gap: The Unmet Civil Legal Needs of Low-income Americans*. Technical Report. Legal Services Corporation.
- [102] Ian René Solano-Kamaiko, Melissa Tan, Joy Ming, Ariel C Avgar, Aditya Vashistha, Madeline Sterling, and Nicola Dell. 2025. "Who is running it?" Towards Equitable AI Deployment in Home Care Work. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*. 1–19.
- [103] William J Stuntz. 1997. The uneasy relationship between criminal procedure and criminal justice. *Yale LJ* 107 (1997), 1.
- [104] Iddo Tavory and Stefan Timmermans. 2014. *Abductive analysis: Theorizing qualitative research*. University of Chicago press.
- [105] Songül Tolan, Marius Miron, Emilia Gómez, and Carlos Castillo. 2019. Why machine learning may lead to unfairness: Evidence from risk assessment for juvenile justice in catalonia. In *Proceedings of the seventeenth international conference on artificial intelligence and law*. 83–92.
- [106] Shengbang Tong, Zhuang Liu, Yuexiang Zhai, Yi Ma, Yann LeCun, and Saining Xie. 2024. Eyes wide shut? exploring the visual shortcomings of multimodal llms. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 9568–9578.
- [107] Smith v. Maryland. 1979. 442 U.S. 735.
- [108] Michelle Vaccaro, Abdullah Almaatouq, and Thomas Malone. 2024. When combinations of humans and AI are useful: A systematic review and meta-analysis. *Nature Human Behaviour* 8, 12 (2024), 2293–2303.
- [109] Zhiyuan Wang, Runze Yan, Sherilyn Francis, Carmen Diaz, Tabor Flickinger, Yufen Lin, Xiao Hu, Laura E Barnes, and Virginia LeBaron. 2025. Stakeholder-centric participation in large language models enhanced health systems. *npj Health Systems* 2, 1 (2025), 22.
- [110] Atharva Rajiv Weginwar, Raghunandan Nuggehalli Ramesha, Jayaprakash Ramsaran, and Suyash Bhogawar. 2025. Mitigating Data Leakage Risks: A Comprehensive Approach to Data Governance in Age of AI. *International Journal of Science and Research (IJSR)* (2025). <https://api.semanticscholar.org/CorpusID:278096880>
- [111] Paul B Wice. 2005. *Public defenders and the American justice system*. Bloomsbury Publishing USA.

- [112] Glen Wilkerson. 1972. Public defenders as their clients see them. *Am. J. Crim. L.* 1 (1972), 141.
- [113] Robert Wolfe, Yiwei Yang, Bill Howe, and Aylin Caliskan. 2023. Contrastive language-vision ai models pretrained on web-scraped multimodal data exhibit sexual objectification bias. In *Proceedings of the 2023 ACM Conference on Fairness, Accountability, and Transparency*. 1174–1185.
- [114] Xiaodong Wu, Ran Duan, and Jianbing Ni. 2024. Unveiling security, privacy, and ethical concerns of ChatGPT. *Journal of information and intelligence* 2, 2 (2024), 102–115.
- [115] Boqiang Zhang, Kehan Li, Zesen Cheng, Zhiqiang Hu, Yuqian Yuan, Guanzheng Chen, Sicong Leng, Yuming Jiang, Hang Zhang, Xin Li, et al. 2025. Videollama 3: Frontier multimodal foundation models for image and video understanding. *arXiv preprint arXiv:2501.13106* (2025).
- [116] Lucia Zheng, Neel Guha, Brandon R. Anderson, Peter Henderson, and Daniel E. Ho. 2021. When Does Pretraining Help? Assessing Self-Supervised Learning for Law and the CaseHOLD Dataset. arXiv:2104.08671 [cs.CL] <https://arxiv.org/abs/2104.08671>
- [117] Yang Zhong and Diane Litman. 2022. Computing and exploiting document structure to improve unsupervised extractive summarization of legal case decisions. *arXiv preprint arXiv:2211.03229* (2022).

A Interview participants

Participant	Affiliation / Current Role	Years of Experience
P1*	Federal public defender	40+
P2*	Federal court-appointed private attorney	7
P3*	Federal investigator	30+
P4*	Federal public defender	10
P5	State public defender	20+
P6	Federal public defender	10
P7	Federal court-appointed private attorney	10+
P8	State public defender	5
P9	Law professor, former public defender	30+
P10	Law professor, former public defender	2
P11	Retired law professor, former public defender	30+
P12	Law professor, former public defender	3
P13	Law professor, former public defender	6
P14	Law professor, former public defender	8

Table 2. Description of participants.

*P1 and P2 were interviewed jointly, and P3 and P4 were interviewed jointly.

B Legal Scenario

Case Example

Initial Stop: Officer Rodriguez observed a 2012 Honda Civic with a broken right tail light traveling northbound on Main Street in Bellevue, Washington. When Officer Rodriguez activated emergency lights to initiate a traffic stop, the vehicle initially slowed but then accelerated.

Pursuit & Search: A pursuit began lasting approximately 5 minutes before the vehicle pulled into a gas station parking lot. Officer Rodriguez approached the car and claimed to detect the odor of marijuana. Based on this and the brief chase, officers conducted a warrantless search of the vehicle.

Evidence Found: 50 grams of methamphetamine divided into 12 small plastic bags; Digital scale with residue; \$780 in cash in various denominations; Cell phone with text messages allegedly discussing drug transactions.

Defendant's Claims: The drugs belonged to a friend who borrowed his car earlier that day; Body camera footage has some audio gaps during crucial moments of the search; Recently attempted to enter a substance abuse program but was placed on a waiting list due to insurance issues.

C Iterative Evolution of Interview Protocol

Our interview approach evolved from broad exploratory questions to targeted theoretical probes as emerging themes guided protocol refinement. This evolution exemplifies grounded theory's commitment to letting theory emerge from data through constant comparative analysis.

C.1 Initial Interview Question Set

Our interviews began with a semi-structured protocol organized around the following core areas:

Prior Experience with AI Tools (10-15 minutes).

- How often, if at all, do you use AI tools inside and outside of work? Could you share what tools you have used?
- What do you think of using AI for work based on your prior experience and what do you know about these AI tools? What do you find promising or frustrating about these tools? What are you most concerned about regarding using AI tools?
- Which part of your work do you value the most? What are the parts that are time consuming but not as rewarding?

Case Deliberation (10 minutes).

- Can you walk us through how you would approach this case?
- Which step is usually the most time consuming or most difficult?

AI-assisted Information Retrieval (15 minutes).

- What service would you use first, and what search would you run?
- Follow-up questions based on participant responses and tool selections.

Thoughts on General AI Tools (10 minutes).

- Which parts of your work do you think is most likely to be automated, and would be most valuable to have AI assistance?
- What is missing in current AI tools?
- Which of these tools would be the most helpful for your work and how would you use them?

Closing Questions (5 minutes).

- Open-ended opportunity for participants to ask questions about AI, the research project, or comment on topics not covered in sufficient detail.

C.2 Evolved Comprehensive Question Set

Through iterative refinement guided by emerging themes, our protocol expanded to include the following questions. However, these questions served as flexible guideposts rather than a rigid script—actual interviews varied considerably based on individual participants’ experiences, interests, and responses, with new questions emerging organically during conversations and others being omitted when not relevant to particular participants’ contexts.

Workflow Mapping & Task Prioritization.

- Looking at these five categories (Defense Strategy, Client Communication, Legal Research, Evidence Investigation, Courtroom Representation), do these five categories representative of your typical tasks? Do you want to edit anything?
- Where do you feel most overwhelmed in your caseload?
- What aspects of your work do you think we don’t fully understand?

Case Strategy & Decision-Making Process.

- (Presenting the prepared case) When you get a new case like this case, what’s your very first step?
- (Presenting the five categories of work) How would you allocate your time across five categories of your task?

- Can you rank them from most to least time-consuming? What is usually the most challenging or difficult?
- How do you prioritize which issues to investigate first?
- What factors most influence your strategic decisions? How do you decide whether to pursue suppression vs. plea bargaining vs. trial?

AI Integration Boundaries & Professional Identity.

- Even if AI could help with various tasks, what work would you always want to do yourself? What aspects of your work feel too important to your role as an advocate to hand over to technology?
- What institutional barriers might prevent AI adoption in public defender offices?
- Do you think attitudes toward AI vary depending on roles (trial v. research attorneys), seniority (entry-level v. management), and employment types (private v. public)? What other factors could influence attitudes toward AI?
- What features would an ideal AI legal assistant have? What would most improve your ability to serve clients effectively?

Technology Assessment & Tool Evaluation. (Optional: when participants mentioned a specific tool)

- Is the output generated by ChatGPT and Westlaw AI helpful? What's missing or concerning?
- What improvements would make this more useful for your practice?
- What would make you more likely to adopt AI tools in your practice?
- What would you want AI developers to understand about public defense work?

Received 20 February 2024; revised 12 March 2024; accepted 5 June 2024